

Machine Learning Homework 6

Machine Learning Homework 6

Environment

Code with detailed explanations

Libraries

Visualization

Part1

Part2

Part3

Part4

Experiments settings and results & Discussion

Part1

Part2

Part3

Part4

Observations and discussion

Environment

- Language: C++
- Standard: C++20

Code with detailed explanations

Libraries

Used library => Corresponding library in homework description

- Eigen => numpy
- OpenCV, matplotlib++ => visualization
- Boost => filesystem, sort
- OpenMP (for parallel acceleration)

Visualization

- **Main part**

Use `fittingHistory` which records the labels for every iterations to make a video.

The duration of one iteration is 1s.

`cv::addWeighted` function is used to combine the clustering results (segmentation mask) with original image.

`drawMask` function is used to convert the labels to segmentation mask.

```

101     kernelKMeans.fit(kernel);
102     const std::vector<Eigen::VectorXi> &fittingHistory = kernelKMeans.getFittingHistory();
103
104     writer.open(imageFile + "_video.mp4", codec, fps, image.size());
105     // check if we succeeded
106     if (!writer.isOpened())
107     {
108         throw std::runtime_error("Could not open the output video file for write");
109     }
110
111     for (int i = 0; i < fps; i++)
112     {
113         writer.write(image);
114     }
115
116     cv::Mat result;
117     cv::Mat mask;
118     for (std::size_t i = 1; i < fittingHistory.size(); i++)
119     {
120         mask = drawMask(fittingHistory[i], numberOfClusters, image.cols, image.rows);
121         cv::addWeighted(image, 0.5, mask, 0.5, 0, result);
122         for (int i = 0; i < fps; i++)
123         {
124             writer.write(result);
125         }
126     }
127
128     writer.release();
129
130     cv::imwrite(imageFile + "_mask.png", mask);
131     cv::imwrite(imageFile + "_final.png", result);

```

- **Draw mask**

Convert the labels to segmentation mask.

`labels` contains the clustering results, from 0 to k-1.

`cv::applyColorMap` is used to map the labels to the corresponding color to avoid duplicate colors.

```

65 cv::Mat drawMask(const Eigen::VectorXi &labels, int numberOfClusters, unsigned int width, unsigned int height)
66 {
67     Eigen::VectorX<unsigned char> maskData = labels.cast<unsigned char>();
68     cv::Mat mask = cv::Mat(cv::Size(width, height), CV_8UC1, reinterpret_cast<void*>(maskData.data()));
69
70     mask *= (255 / numberOfClusters);
71
72     cv::Mat bgrMask;
73     cv::cvtColor(mask, bgrMask, cv::COLOR_GRAY2BGR);
74     cv::applyColorMap(bgrMask, bgrMask, cv::COLORMAP_COOL);
75     return bgrMask;
76 }

```

Part1

- **Kernel K-Means**

- Pseudo-code

Ignore weight

ALGORITHM 1: Basic Batch Weighted Kernel k -means.

KERNEL_KMEANS_BATCH($K, k, w, t_{max}, \{\pi_c^{(0)}\}_{c=1}^k, \{\pi_c\}_{c=1}^k$)

Input: K : kernel matrix, k : number of clusters, w : weights for each point, t_{max} : optional maximum number of iterations, $\{\pi_c^{(0)}\}_{c=1}^k$: optional initial clusters

Output: $\{\pi_c\}_{c=1}^k$: final partitioning of the points

1. If no initial clustering is given, initialize the k clusters $\pi_1^{(0)}, \dots, \pi_k^{(0)}$ (i.e., randomly). Set $t = 0$.
2. For each point \mathbf{a}_i and every cluster c , compute

$$d(\mathbf{a}_i, \mathbf{m}_c) = K_{ii} - \frac{2 \sum_{\mathbf{a}_j \in \pi_c^{(t)}} w_j K_{ij}}{\sum_{\mathbf{a}_j \in \pi_c^{(t)}} w_j} + \frac{\sum_{\mathbf{a}_j, \mathbf{a}_l \in \pi_c^{(t)}} w_j w_l K_{jl}}{(\sum_{\mathbf{a}_j \in \pi_c^{(t)}} w_j)^2}.$$

3. Find $c^*(\mathbf{a}_i) = \operatorname{argmin}_c d(\mathbf{a}_i, \mathbf{m}_c)$, resolving ties arbitrarily. Compute the updated clusters as

$$\pi_c^{(t+1)} = \{\mathbf{a} : c^*(\mathbf{a}) = c\}.$$

4. If not converged or $t_{max} > t$, set $t = t + 1$ and go to Step 2; Otherwise, stop and output final clusters $\{\pi_c^{(t+1)}\}_{c=1}^k$.

◦ Main function

Arguments

path: the image data folder

numberOfClusters: the number of clusters

init: the selected initialization method, random or k-means++

gamma1, gamma2: the hyper-parameter of RBF kernel

```

78 void run(const fs::path &path, int numberOfClusters, mlhw6::KMeansInitMethods init, double gamma1, double gamma2)
79 {
80     int fps = 30;
81     int codec = cv::VideoWriter::fourcc('m', 'p', '4', 'v');
82     cv::VideoWriter writer;
83
84     mlhw6::KernelKMeans kernelKMeans(numberOfClusters, 200, 1234, init);
85
86     for (auto imageFile : IMAGE_FILES)
87     {
88         std::cout << imageFile << std::endl;
89
90         auto fileName = (path / imageFile).generic_string();
91
92         // read image
93         auto image = cv::imread(fileName, cv::ImreadModes::IMREAD_COLOR);
94
95         // extract RGB values and coordinates
96         auto [pixels, coordinates] = preprocess(image);
97
98         // calculate kernel
99         auto kernel = calculateKernel(pixels, coordinates, gamma1, gamma2);
100
101         kernelKMeans.fit(kernel);
102         const std::vector<Eigen::VectorXi> &fittingHistory = kernelKMeans.getFittingHistory();
103
104         writer.open(imageFile + "_video.mp4", codec, fps, image.size());

```

◦ kmeans.h header: k-means related classes

Follow the scikit-learn logic design.

```

7   enum KMeansInitMethods
8   {
9       Random,
10      Kmeansplusplus,
11  };
12
13  class BaseKMeans
14  {
15  public:
16      BaseKMeans(int numberOfClusters, KMeansInitMethods init = KMeansInitMethods::Kmeansplusplus, int maximumEpochs = 200, int seed = 1234);
17
18      virtual void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) = 0;
19      virtual Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const = 0;
20      virtual Eigen::VectorXi fitAndPredict(const Eigen::Ref<const Eigen::MatrixXd> &x);
21
22      const std::vector<Eigen::VectorXi> &getFittingHistory() const;
23
24  protected:
25      Eigen::VectorXi initializeCenters(const Eigen::Ref<const Eigen::MatrixXd> &x, KMeansInitMethods init, int seed, bool precomputed = false) const;
26
27      int numberOfClusters;
28      int maximumEpochs;
29      int seed;
30      KMeansInitMethods init;
31
32      std::vector<Eigen::VectorXi> fittingHistory;
33  };
34
35  class KMeans : public virtual BaseKMeans
36  {
37  public:
38      using BaseKMeans::BaseKMeans;
39
40      void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
41      Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const override;
42
43  private:
44      Eigen::VectorXi assignLabels(const Eigen::Ref<const Eigen::MatrixXd> &x) const;
45      Eigen::MatrixXd centers;
46  };
47
48  class KernelKMeans : public virtual BaseKMeans
49  {
50  public:
51      using BaseKMeans::BaseKMeans;
52
53      void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
54      Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const override;
55
56  private:
57      Eigen::VectorXi assignLabels(const Eigen::Ref<const Eigen::MatrixXd> &x) const;
58  };

```

- `KernelKMeans` constructor

```

174      KernelKMeans::KernelKMeans(int numberOfClusters, int maximumEpochs, int seed, KMeansInitMethods init) : BaseKMeans(numberOfClusters, maximumEpochs, seed, init)
175      {
176      }
177
178  BaseKMeans::BaseKMeans(int numberOfClusters, int maximumEpochs, int seed, KMeansInitMethods init) : numberOfClusters(numberOfClusters), maximumEpochs(maximumEpochs), seed(seed), init(init)
179  {
180      if (!numberOfClusters > 0)
181      {
182          throw std::runtime_error("The number of clusters should be larger than 0.");
183      }
184  }

```

- `preprocess` function: extract RGB values and coordinates

Arguments

image: (H, W, 3), BGR values

```

25  std::pair<Eigen::MatrixX3d, Eigen::MatrixX2i> preprocess(const cv::Mat &image)
26  {
27      auto rows = image.rows;
28      auto columns = image.cols;
29      auto size = rows * columns;
30
31      cv::Mat rgb;
32      cv::cvtColor(image, rgb, cv::COLOR_BGR2RGB);
33
34      std::vector<int> coordinates(size * 2);
35      #pragma omp parallel for collapse(2)
36      for (unsigned int i = 0; i < rows; i++)
37      {
38          for (unsigned int j = 0; j < columns; j++)
39          {
40              auto index = (i * columns + j) * 2;
41              coordinates[index] = i;
42              coordinates[index + 1] = j;
43          }
44      }
45
46      using MatrixX3ucRowMajor = Eigen::Matrix<unsigned char, Eigen::Dynamic, 3, Eigen::RowMajor>;
47      using MatrixX2iRowMajor = Eigen::Matrix<int, Eigen::Dynamic, 2, Eigen::RowMajor>;
48
49      return std::make_pair<Eigen::MatrixX3d, Eigen::MatrixX2i>(
50          MatrixX3ucRowMajor::Map(rgb.data, size, 3).cast<double>(), MatrixX2iRowMajor::Map(coordinates.data(), size, 2));
51  }

```

- `calculatekernel` function: calculate all kernel values

Arguments

pixels: (N, 3), RGB values

coordinates: (N, 2), coordinates

gamma1: γ_c scalar

gamma2: γ_s scalar

Formula

$$k(x, x') = e^{-\gamma_s \|S(x) - S(x')\|^2} \times e^{-\gamma_c \|C(x) - C(x')\|^2}$$

$S(x)$ is the spatial information (coordinate)

$C(x)$ is the color information (RGB)

```
55 Eigen::MatrixXd calculateKernel(const Eigen::MatrixX3d &pixels, const Eigen::MatrixX2i &coordinates, double gamma1, double gamma2)
56 {
57     auto colorKernel = mlhw6::rbf(pixels, pixels, gamma1);
58
59     const Eigen::MatrixX2d &tmp = coordinates.cast<double>();
60     auto coordinateKernel = mlhw6::rbf(tmp, tmp, gamma2);
61
62     return colorKernel.cwiseProduct(coordinateKernel);
63 }
```

- o `rbf` function: RBF kernel

Arguments

x1: x vector

x2: x' vector

gamma: γ scalar

Formula

$$k(x, x') = e^{-\gamma \|x - x'\|^2}$$

```
9 template <typename DerivedA, typename DerivedB, typename Out = Eigen::Matrix<double, DerivedA::RowsAtCompileTime, DerivedB::RowsAtCompileTime>>
10 Out rbf(const Eigen::MatrixBase<DerivedA> &x1, const Eigen::MatrixBase<DerivedB> &x2, double gamma)
11 {
12     Out result(x1.rows(), x2.rows());
13     #pragma omp parallel for
14     for (Eigen::Index i = 0; i < x1.rows(); i++)
15     {
16         result.row(i) = (-gamma * (x2.rowwise() - x1.row(i)).rowwise().squaredNorm().transpose()).array().exp();
17     }
18     return result;
19 }
```

- o `kernelKMeans.fit` function: fit the data

Previously, we already calculated the kernel.

So, we use the precomputed kernel values directly.

Arguments

x: (N, N), the kernel values (gram matrix, similarity matrix)

Steps

1. Pick k centers
2. Calculate the cost between the data and centers
3. Assign the label which has the smallest distance to the data
4. Keep repeating 2, 3 step until the labels are not changed

`fittingHistory` is used to store the labels for every iterations.

Note: At the line 183, we do $1 - x$ to get the distance matrix.

```
178 void KernelKMeans::fit(const Eigen::Ref<const Eigen::MatrixXd> &x)
179 {
180     this->fittingHistory = std::vector<Eigen::VectorXi>{Eigen::VectorXi::Constant(x.rows(), -1)};
181
182     // x is similarity matrix (gram matrix)
183     auto centers = this->initializeCenters(1 - x.array(), this->init, this->seed, true);
184     this->fittingHistory.back()(centers).setLinSpaced(0, this->numberOfClusters - 1);
185
186     int epoch = 0;
187     bool sameLabels = false;
188     do
189     {
190         this->fittingHistory.push_back(this->assignLabels(x));
191
192         sameLabels = (this->fittingHistory[epoch].array() == this->fittingHistory[epoch + 1].array()).all();
193         epoch++;
194     } while (!sameLabels && epoch < this->maximumEpochs);
195
196     std::cout << "Finished at " << epoch << " epoch" << std::endl;
197 }
```

- `assignLabels` function: calculate the cost and assign labels

Arguments

`x`: (N, N), the kernel values

Steps

1. Calculate the cost between the data and centers

$$\mathbf{k}(x_j, x_j) - \frac{2}{|C_k|} \sum_n \alpha_{kn} \mathbf{k}(x_j, x_n) + \frac{1}{|C_k|^2} \sum_p \sum_q \alpha_{kp} \alpha_{kq} \mathbf{k}(x_p, x_q)$$

2. Assign the label which has the smallest distance to the data

```

204 Eigen::VectorXi KernelKMeans::assignLabels(const Eigen::Ref<const Eigen::MatrixXd> x) const
205 {
206     Eigen::MatrixXd distance(x.rows(), this->numberOfClusters);
207     #pragma omp parallel for
208     for (int k = 0; k < this->numberOfClusters; k++)
209     {
210         Eigen::VectorXd selector = this->fittingHistory.back().cwiseEqual(k).cast<double>();
211         auto numberOfXInKCluster = selector.sum();
212         Eigen::MatrixXd xToKCluster = x * selector.asDiagonal();
213
214         Eigen::VectorXd secondTerm = 2 * xToKCluster.rowwise().sum() / numberOfXInKCluster;
215         auto thirdTerm = (selector.asDiagonal() * xToKCluster).sum() / std::pow(numberOfXInKCluster, 2);
216         distance(Eigen::all, k) = (x.diagonal() - secondTerm).array() + thirdTerm;
217     }
218
219     std::vector<int> labels(x.rows());
220     #pragma omp parallel for
221     for (Eigen::Index i = 0; i < x.rows(); i++)
222     {
223         // find nearest neighbor
224         distance.row(i).minCoeff(&labels[i]);
225     }
226     return Eigen::VectorXi::Map(labels.data(), labels.size());
227 }

```

• Spectral Clustering

- Pseudo-code
 - Spectral Clustering (ratio cut)

Unnormalized spectral clustering

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct.

- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the unnormalized Laplacian L .
- Compute the first k eigenvectors u_1, \dots, u_k of L .
- Let $U \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors u_1, \dots, u_k as columns.
- For $i = 1, \dots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i -th row of U .
- Cluster the points $(y_i)_{i=1, \dots, n}$ in \mathbb{R}^k with the k -means algorithm into clusters C_1, \dots, C_k .

Output: Clusters A_1, \dots, A_k with $A_i = \{j \mid y_j \in C_i\}$.

- Normalized Spectral Clustering (normalized cut)

Normalized spectral clustering according to Shi and Malik (2000)

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct.

- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the unnormalized Laplacian L .
- Compute the first k generalized eigenvectors u_1, \dots, u_k of the generalized eigenproblem $Lu = \lambda Du$.
- Let $U \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors u_1, \dots, u_k as columns.
- For $i = 1, \dots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i -th row of U .
- Cluster the points $(y_i)_{i=1, \dots, n}$ in \mathbb{R}^k with the k -means algorithm into clusters C_1, \dots, C_k .

Output: Clusters A_1, \dots, A_k with $A_i = \{j \mid y_j \in C_i\}$.

- Main function

Arguments

path: the image data folder

model: Spectral Clustering instance

numberOfClusters: the number of clusters

gamma1, gamma2: the hyper-parameter of RBF kernel

```
switch (type)
{
case SpectralClusteringType::Unnormalized:
{
    auto model = mlhw6::SpectralClustering(numberOfClusters, init);
    run(path, model, numberOfClusters, gamma1, gamma2);
    break;
}
case SpectralClusteringType::Normalized:
{
    auto model = mlhw6::NormalizedSpectralClustering(numberOfClusters, init);
    run(path, model, numberOfClusters, gamma1, gamma2);
    break;
}
default:
    std::cerr << "Unknown spectral clustering type." << std::endl;
    return 1;
}
```

```
85 void run(const fs::path &path, mlhw6::BaseSpectralClustering &model, int numberOfClusters, double gamma1, double gamma2)
86 {
87     int fps = 30;
88     int codec = cv::VideoWriter::fourcc('m', 'p', '4', 'v');
89     cv::VideoWriter writer;
90
91     for (auto imageFile : IMAGE_FILES)
92     {
93         std::cout << imageFile << std::endl;
94
95         auto fileName = (path / imageFile).generic_string();
96
97         // read image
98         auto image = cv::imread(fileName, cv::ImreadModes::IMREAD_COLOR);
99
100        // extract RGB values and coordinates
101        auto [pixels, coordinates] = preprocess(image);
102
103        // calculate kernel
104        auto kernel = calculateKernel(pixels, coordinates, gamma1, gamma2);
105
106        model.fit(kernel);
107        const std::vector<Eigen::VectorXi> &fittingHistory = model.getFittingHistory();
108
109        writer.open(imageFile + "_video.mp4", codec, fps, image.size());
```

- `spectral.h` header: spectral clustering related classes

Follow the scikit-learn logic design.

```
9 class BaseSpectralClustering
10 {
11 public:
12     BaseSpectralClustering(int numberOfClusters, KMeansInitMethods init = KMeansInitMethods::Kmeansplusplus, int maximumEpochs = 200, int seed = 1234);
13
14     virtual void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) = 0;
15     virtual Eigen::VectorXi fitAndPredict(const Eigen::Ref<const Eigen::MatrixXd> &x);
16
17     const std::vector<Eigen::VectorXi> &getFittingHistory() const;
18     const Eigen::MatrixXd &getEigenMatrix() const;
19
20 protected:
21     KMeans kMeans;
22     Eigen::MatrixXd eigenMatrix;
23
24     int numberOfClusters;
25     unsigned int numberOfThreads;
26 };
27
28 class SpectralClustering : public virtual BaseSpectralClustering
29 {
30 public:
31     using BaseSpectralClustering::BaseSpectralClustering;
32
33     void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
34 };
35
36 class NormalizedSpectralClustering : public virtual BaseSpectralClustering
37 {
38 public:
39     using BaseSpectralClustering::BaseSpectralClustering;
40
41     void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
42 };
```

- `kmeans.h` header: k-means related classes

Follow the scikit-learn logic design.

```

7   enum KMeansInitMethods
8   {
9       Random,
10      Kmeansplusplus,
11  };
12
13  class BaseKMeans
14  {
15  public:
16      BaseKMeans(int numberOfClusters, KMeansInitMethods init = KMeansInitMethods::Kmeansplusplus, int maximumEpochs = 200, int seed = 1234);
17
18      virtual void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) = 0;
19      virtual Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const = 0;
20      virtual Eigen::VectorXi fitAndPredict(const Eigen::Ref<const Eigen::MatrixXd> &x);
21
22      const std::vector<Eigen::VectorXi> &getFittingHistory() const;
23
24  protected:
25      Eigen::VectorXi initializeCenters(const Eigen::Ref<const Eigen::MatrixXd> &x, KMeansInitMethods init, int seed, bool precomputed = false) const;
26
27      int numberOfClusters;
28      int maximumEpochs;
29      int seed;
30      KMeansInitMethods init;
31
32      std::vector<Eigen::VectorXi> fittingHistory;
33  };
34
35  class KMeans : public virtual BaseKMeans
36  {
37  public:
38      using BaseKMeans::BaseKMeans;
39
40      void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
41      Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const override;
42
43  private:
44      Eigen::VectorXi assignLabels(const Eigen::Ref<const Eigen::MatrixXd> &x) const;
45      Eigen::MatrixXd centers;
46  };
47
48  class KernelKMeans : public virtual BaseKMeans
49  {
50  public:
51      using BaseKMeans::BaseKMeans;
52
53      void fit(const Eigen::Ref<const Eigen::MatrixXd> &x) override;
54      Eigen::VectorXi predict(const Eigen::Ref<const Eigen::MatrixXd> &x) const override;
55
56  private:
57      Eigen::VectorXi assignLabels(const Eigen::Ref<const Eigen::MatrixXd> &x) const;
58  };

```

- `preprocess` function: extract RGB values and coordinates

Arguments

image: (H, W, 3), BGR values

```

25  std::pair<Eigen::MatrixX3d, Eigen::MatrixX2i> preprocess(const cv::Mat &image)
26  {
27      auto rows = image.rows;
28      auto columns = image.cols;
29      auto size = rows * columns;
30
31      cv::Mat rgb;
32      cv::cvtColor(image, rgb, cv::COLOR_BGR2RGB);
33
34      std::vector<int> coordinates(size * 2);
35      #pragma omp parallel for collapse(2)
36      for (unsigned int i = 0; i < rows; i++)
37      {
38          for (unsigned int j = 0; j < columns; j++)
39          {
40              auto index = (i * columns + j) * 2;
41              coordinates[index] = i;
42              coordinates[index + 1] = j;
43          }
44      }
45
46      using MatrixX3ucRowMajor = Eigen::Matrix<unsigned char, Eigen::Dynamic, 3, Eigen::RowMajor>;
47      using MatrixX2iRowMajor = Eigen::Matrix<int, Eigen::Dynamic, 2, Eigen::RowMajor>;
48
49      return std::make_pair<Eigen::MatrixX3d, Eigen::MatrixX2i>(
50          MatrixX3ucRowMajor::Map(rgb.data, size, 3).cast<double>(), MatrixX2iRowMajor::Map(coordinates.data(), size, 2));
51  }

```

- `calculatekernel` function: calculate all kernel values

Arguments

pixels: (N, 3), RGB values

coordinates: (N, 2), coordinates

gamma1: γ_c scalar

gamma2: γ_s scalar

Formula

$$k(x, x') = e^{-\gamma_s \|S(x) - S(x')\|^2} \times e^{-\gamma_c \|C(x) - C(x')\|^2}$$

$S(x)$ is the spatial information (coordinate)

$C(x)$ is the color information (RGB)


```

55 Eigen::MatrixXd calculateKernel(const Eigen::MatrixX3d &pixels, const Eigen::MatrixX2i &coordinates, double gamma1, double gamma2)
56 {
57     auto colorKernel = mlhw6::rbf(pixels, pixels, gamma1);
58
59     const Eigen::MatrixX2d &tmp = coordinates.cast<double>();
60     auto coordinateKernel = mlhw6::rbf(tmp, tmp, gamma2);
61
62     return colorKernel.cwiseProduct(coordinateKernel);
63 }

```

- **rbf** function: RBF kernel

Arguments

x_1 : x vector

x_2 : x' vector

γ : scalar

Formula

$$k(x, x') = e^{-\gamma \|x - x'\|^2}$$

```

9     template <typename DerivedA, typename DerivedB, typename Out = Eigen::Matrix<double, DerivedA::RowsAtCompileTime, DerivedB::RowsAtCompileTime>>
10     Out rbf(const Eigen::MatrixBase<DerivedA> &x1, const Eigen::MatrixBase<DerivedB> &x2, double gamma)
11     {
12         Out result(x1.rows(), x2.rows());
13         #pragma omp parallel for
14         for (Eigen::Index i = 0; i < x1.rows(); i++)
15         {
16             result.row(i) = (-gamma * (x2.rowwise() - x1.row(i)).rowwise().squaredNorm().transpose()).array().exp();
17         }
18         return result;
19     }

```

- **model_fit** function: fit the data

Previously, we already calculated the kernel.

So, we use the precomputed kernel values directly.

Arguments

x : (N, N), the kernel values (gram matrix, similarity matrix)

- Spectral Clustering (ratio cut)

Algorithm

Unnormalized spectral clustering

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct.

- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the unnormalized Laplacian L .
- Compute the first k eigenvectors u_1, \dots, u_k of L .
- Let $U \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors u_1, \dots, u_k as columns.
- For $i = 1, \dots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i -th row of U .
- Cluster the points $(y_i)_{i=1, \dots, n}$ in \mathbb{R}^k with the k -means algorithm into clusters C_1, \dots, C_k .

Output: Clusters A_1, \dots, A_k with $A_i = \{j \mid y_j \in C_i\}$.

```

48 void SpectralClustering::fit(const Eigen::Ref<const Eigen::MatrixXd> &x)
49 {
50     Eigen::DiagonalMatrix<double, Eigen::Dynamic> degreeMatrix = x.rowwise().sum().asDiagonal();
51
52     // L = D - W
53     Eigen::MatrixXd laplacianMatrix = -x;
54     laplacianMatrix.diagonal() += degreeMatrix.diagonal();
55
56     // solve eigen decomposition
57     Eigen::EigenSolver<Eigen::MatrixXd> solver(laplacianMatrix, true);
58     const Eigen::VectorXcd &eigenValues = solver.eigenvalues();
59     const Eigen::MatrixXcd &eigenVectors = solver.eigenvectors();
60
61     std::cout << eigenValues.topRows(5) << std::endl;
62     std::cout << eigenVectors.leftCols(5) << std::endl;
63
64     // sort eigenvalues
65     std::vector<std::pair<double, Eigen::Index>> eigenPairs(eigenValues.rows());
66     #pragma omp parallel for
67     for (Eigen::Index i = 0; i < eigenValues.rows(); i++)
68     {
69         eigenPairs[i] = std::make_pair(eigenValues[i].real(), i);
70     }
71     boost::sort::parallel_stable_sort(eigenPairs.begin(), eigenPairs.end(), this->numberOfThreads);
72     // std::partial_sort(eigenPairs.begin(), eigenPairs.begin() + this->numberOfClusters, eigenPairs.end());
73
74     // pick k eigenvectors
75     this->eigenMatrix = Eigen::MatrixXd(eigenVectors.rows(), this->numberOfClusters);
76     #pragma omp parallel for
77     for (int i = 0; i < this->numberOfClusters; i++)
78     {
79         this->eigenMatrix.col(i) = eigenVectors.col(eigenPairs[i].second).real();
80     }
81
82     // perform k-means clustering on eigen space
83     this->kMeans.fit(this->eigenMatrix);
84 }

```

■ Normalized Spectral Clustering (normalized cut)

Algorithm

Normalized spectral clustering according to Shi and Malik (2000)

Input: Similarity matrix $S \in \mathbb{R}^{n \times n}$, number k of clusters to construct.

- Construct a similarity graph by one of the ways described in Section 2. Let W be its weighted adjacency matrix.
- Compute the unnormalized Laplacian L .
- Compute the first k generalized eigenvectors u_1, \dots, u_k of the generalized eigenproblem $Lu = \lambda Du$.
- Let $U \in \mathbb{R}^{n \times k}$ be the matrix containing the vectors u_1, \dots, u_k as columns.
- For $i = 1, \dots, n$, let $y_i \in \mathbb{R}^k$ be the vector corresponding to the i -th row of U .
- Cluster the points $(y_i)_{i=1, \dots, n}$ in \mathbb{R}^k with the k -means algorithm into clusters C_1, \dots, C_k .

Output: Clusters A_1, \dots, A_k with $A_i = \{j | y_j \in C_i\}$.

```

89 void NormalizedSpectralClustering::fit(const Eigen::Ref<const Eigen::MatrixXd> &x)
90 {
91     Eigen::DiagonalMatrix<double, Eigen::Dynamic> degreeMatrix = x.rowwise().sum().asDiagonal();
92
93     // L = D - W
94     Eigen::MatrixXd laplacianMatrix = -x;
95     laplacianMatrix.diagonal() += degreeMatrix.diagonal();
96
97     // solve generalized eigen decomposition
98     Eigen::GeneralizedEigenSolver<Eigen::MatrixXd> solver(laplacianMatrix, degreeMatrix, true);
99     const Eigen::VectorXcd &eigenValues = solver.eigenvalues();
100     const Eigen::MatrixXcd &eigenVectors = solver.eigenvectors();
101
102     std::cout << eigenValues.topRows(5) << std::endl;
103     std::cout << eigenVectors.leftCols(5) << std::endl;
104
105     // sort eigenvalues
106     std::vector<std::pair<double, Eigen::Index>> eigenPairs(eigenValues.rows());
107     #pragma omp parallel for
108     for (Eigen::Index i = 0; i < eigenValues.rows(); i++)
109     {
110         eigenPairs[i] = std::make_pair(eigenValues[i].real(), i);
111     }
112     boost::sort::parallel_stable_sort(eigenPairs.begin(), eigenPairs.end(), this->numberOfThreads);
113     // std::partial_sort(eigenPairs.begin(), eigenPairs.begin() + this->numberOfClusters, eigenPairs.end());
114
115     // pick k eigenvectors
116     this->eigenMatrix = Eigen::MatrixXd(eigenVectors.rows(), this->numberOfClusters);
117     #pragma omp parallel for
118     for (int i = 0; i < this->numberOfClusters; i++)
119     {
120         this->eigenMatrix.col(i) = eigenVectors.col(eigenPairs[i].second).real();
121     }
122
123     // perform k-means clustering on eigen space
124     this->kMeans.fit(this->eigenMatrix);
125 }

```

- `kmeans_fit` function: fit the data

Use the points in eigen space to perform k-means clustering.

Steps

1. Pick k centers

2. E step

1. Calculate the cost between the data and centers
2. Assign the label which has the smallest distance to the data

3. M step

Recalculate the centers by averaging the points which belong to the same cluster.

4. Keep repeating 2, 3 step until the labels are not changed

```
123 void KMeans::fit(const Eigen::Ref<const Eigen::MatrixXd> &x)
124 {
125     this->fittingHistory = std::vector<Eigen::VectorXi>{Eigen::VectorXi::Constant(x.rows(), -1)};
126
127     auto centers = this->initializeCenters(x, this->init, this->seed);
128     this->fittingHistory.back()(centers).setLinSpaced(0, this->numberOfClusters - 1);
129     this->centers = x(centers, Eigen::all);
130
131     int epoch = 0;
132     bool sameLabels = false;
133     do
134     {
135         std::cout << "Epoch: " << epoch << std::endl;
136
137         // E step
138         this->fittingHistory.push_back(this->assignLabels(x));
139
140         // M step
141         #pragma omp parallel for
142         for (int k = 0; k < this->numberOfClusters; k++)
143         {
144             Eigen::VectorXd selector = this->fittingHistory.back().cwiseEqual(k).cast<double>();
145             // calculate center
146             this->centers.row(k) = (selector.asDiagonal() * x).colwise().sum() / selector.sum();
147         }
148
149         sameLabels = (this->fittingHistory[epoch].array() == this->fittingHistory[epoch + 1].array()).all();
150         epoch++;
151     } while (!sameLabels && epoch < this->maximumEpochs);
152 }
```

- o `assignLabels` function: calculate the cost and assign labels

Arguments

x: (N, N), the kernel values

Steps

1. Calculate the cost (Euclidean distance) between the data and centers
2. Assign the label which has the smallest distance to the data

```
159 Eigen::VectorXi KMeans::assignLabels(const Eigen::Ref<const Eigen::MatrixXd> &x) const
160 {
161     std::vector<int> labels(x.rows());
162     #pragma omp parallel for
163     for (Eigen::Index i = 0; i < x.rows(); i++)
164     {
165         // find nearest neighbor
166         (this->centers.rowwise() - x.row(i)).rowwise().squaredNorm().minCoeff(&labels[i]);
167     }
168     return Eigen::VectorXi::Map(labels.data(), labels.size());
169 }
```

Part2

Use command to control the number of clusters and other parameters.

• Kernel K-Means

init

0 => use random initialization

1 => use k-means++ initialization

```
int main(int argc, char *argv[])
{
    if (argc < 6)
    {
        std::cerr << "Usage: " << argv[0] << " <data path> <number of cluster> <init> <gamma1> <gamma2>" << std::endl;
        return 1;
    }
}
```

- **Spectral Clustering**

spectral clustering type

0 => use spectral clustering algorithm

1 => use normalized clustering algorithm

init

0 => use random initialization

1 => use k-means++ initialization

```
int main(int argc, char *argv[])
{
    if (argc < 7)
    {
        std::cerr << "Usage: " << argv[0] << " <data path> <number of cluster> <spectral clustering type> <init> <gamma1> <gamma2>" << std::endl;
        return 1;
    }
}
```

Part3

- **initializeCenters** function: pick k centers initialized by the selected method

Arguments

x: (N, N) precomputed distance matrix or (N, features) data

init: the selected initialization method, random or k-means++

seed: the random seed

precomputed: x is the precomputed distance matrix or not.

```
Eigen::VectorXi BaseKMeans::initializeCenters(const Eigen::Ref<const Eigen::MatrixXd> &x, KMeansInitMethods init, int seed, bool precomputed) const
{
    if (precomputed)
    {
        if (x.rows() != x.cols())
        {
            throw std::runtime_error("The precomputed x should be a squared matrix.");
        }
    }

    switch (init)
    {
        case KMeansInitMethods::Random:
            return randomInitialization(x, this->numberOfClusters, seed);
        case KMeansInitMethods::Kmeansplusplus:
            return kMeansPlusPlusInitialization(x, this->numberOfClusters, seed, precomputed);
        default:
            throw std::runtime_error("The initialization method is not supported.");
    }
}
```

- **Random initialization**

Arguments

x: (N, N) precomputed distance matrix or (N, features) data

numberOfClusters: k clusters

seed: the random seed

Steps

1. generate the sequence of indexes, 0 ~ N-1
2. shuffle the sequence
3. pick the top k rows as the centers

```

65 Eigen::VectorXi randomInitialization(const Eigen::Ref<const Eigen::MatrixXd> &x, int numberOfClusters, int seed)
66 {
67     auto rng = std::mt19937_64(seed);
68     std::vector<int> sequence(x.rows());
69     std::iota(sequence.begin(), sequence.end(), 0);
70     std::shuffle(sequence.begin(), sequence.end(), rng);
71
72     Eigen::Map<Eigen::VectorXi> tmp = Eigen::VectorXi::Map(sequence.data(), sequence.size());
73     return tmp.topRows(numberOfClusters);
74 }

```

- **K-Means++ initialization**

Arguments

x: (N, N) precomputed distance matrix or (N, features) data

numberOfClusters: k clusters

seed: the random seed

Steps

1. Choose one center uniformly at random among the data points.
2. For each data point x not chosen yet, compute $D(x)^2$ (the squared Euclidean distance) or use the precomputed distance matrix, the distance between x and the nearest center that has already been chosen.
3. Choose one new data point at random as a new center, using a weighted probability distribution where a point x is chosen with probability proportional to $D(x)^2$. (The farthest point will be chosen.)

```

12 Eigen::VectorXi kMeansPlusPlusInitialization(const Eigen::Ref<const Eigen::MatrixXd> &x, int numberOfClusters, int seed, bool precomputed)
13 {
14     auto rng = std::mt19937_64(seed);
15     std::vector<int> candidates;
16
17     // 1. Choose one center uniformly at random among the data points.
18     // closed interval [0, rows - 1]
19     auto sampler = std::uniform_int_distribution(0, static_cast<int>(x.rows()) - 1);
20     candidates.push_back(sampler(rng));
21     numberOfClusters--;
22
23     // 2. For each data point x not chosen yet, compute D(x)^2, the distance between x and the nearest center that has already been chosen.
24     Eigen::MatrixXd distances(x.rows(), x.rows());
25     if (!precomputed)
26     {
27         #pragma omp parallel for
28         for (Eigen::Index i = 0; i < x.rows(); i++)
29         {
30             distances.row(i) = (x.rowwise() - x.row(i)).rowwise().squaredNorm().transpose();
31         }
32     }
33     else
34     {
35         distances = x;
36     }
37
38     auto probabilityDistribution = std::uniform_real_distribution();
39     while (numberOfClusters > 0)
40     {
41         Eigen::Map<Eigen::VectorXi> eigenCandidates = Eigen::VectorXi::Map(candidates.data(), candidates.size());
42
43         // 3. Choose one new data point at random as a new center, using a weighted probability distribution where a point x is chosen with probability proportional to D(x)^2.
44         Eigen::VectorXd weights = distances(Eigen::all, eigenCandidates).rowwise().minCoeff();
45         weights /= weights.sum();
46
47         auto probability = probabilityDistribution(rng);
48         for (int i = 0; i < weights.rows(); i++)
49         {
50             auto weight = weights[i];
51             if (probability < weight)
52             {
53                 candidates.push_back(i);
54                 break;
55             }
56             probability -= weight;
57         }
58
59         numberOfClusters--;
60     }
61 }

```

Part4

Visualize the eigen space.

```

85 void plotEigenSpace(const std::string &path, const Eigen::MatrixXd &eigenMatrix)
86 {
87     matplotlib::figure();
88
89     const Eigen::VectorXd &vecX = eigenMatrix.col(0);
90     const Eigen::VectorXd &vecY = eigenMatrix.col(1);
91
92     std::vector<double> x(vecX.begin(), vecX.end());
93     std::vector<double> y(vecY.begin(), vecY.end());
94     matplotlib::scatter(x, y);
95     matplotlib::save(path);
96 }

```

Experiments settings and results & Discussion

Part1

- **Kernel K-Means**

Number of clusters: 2

Initialization method: random


Gamma1: 0.00001

Gamma2: 0.00001

- image1

- [Video](#)

- Final



- Mask




- image2

- [Video](#)

- Final



- Mask



- Spectral Clustering

Part2

- Kernel K-Means

Initialization method: random


Gamma1: 0.00001

Gamma2: 0.00001

- $K = 3 \cdot \text{image1}$

- [Video](#)

- Final




- Mask



- $K = 3 \cdot \text{image2}$

- [Video](#)



- Final




- Mask

- $K = 4 \cdot \text{image1}$

- [Video](#)
- Final




- Mask




- $K = 4 \cdot \text{image2}$

- [Video](#)
- Final



- Mask



- Spectral Clustering

Part3

- **Kernel K-Means**

I think there is no significant difference between random and k-means++.

Initialization method: k-means++


Gamma1: 0.00001

Gamma2: 0.00001

- K = 2 · image1

- [Video](#)

- Final




- Mask




- K = 2 · image2

- [Video](#)

- Final




- Mask



- $K = 3 \cdot \text{image1}$

- [Video](#)

- Final



- Mask




- $K = 3 \cdot \text{image2}$

- [Video](#)

- Final



- Mask




- $K = 4 \cdot \text{image1}$

- [Video](#)

- Final

- Mask




- $K = 4 \cdot \text{image2}$

- [Video](#)


- Final



- Mask



- Spectral Clustering



Part4

Observations and discussion

- Coming soon...